📋 TRAINER SCRIPT

Week 2 · Session 2 — ETL vs ELT | Batch vs Real-Time | Industry Tools

Total Duration: 60 Minutes | Conceptual Only — Practical in W13

⚙️ PRE-SESSION CHECKLIST

- [ ] PPT loaded and projected
- [ ] Whiteboard/marker available (for live flow diagrams)
- [ ] Quiz printed or sent digitally
- [ ] Learners have done W2S1 (Data Warehousing basics)

🎯 SESSION GOALS

By end of this session, learners should be able to: 1. Explain ETL and ELT with real examples 2. Differentiate batch vs real-time processing 3. Name at least 5 industry tools and map them to use cases

---

📍 SLIDE 1 — Title Slide

⏱ Time: 2 minutes

Say:

> "Good morning/afternoon everyone. Welcome to Week 2, Session 2. Last session we covered Data Warehousing — today we're going one level deeper into HOW data actually gets into those warehouses.
>
> This is one of the most practical and interview-relevant topics in data engineering. By the end of this session, you'll be able to answer questions like: 'How does Swiggy know what to recommend to you in real time?' or 'Why do banks process your salary at midnight rather than instantly?' — all of this is about pipelines."

Trainer Tip: Create a little buzz. Ask: *"Has anyone heard the term ETL before? Put your hand up."* This gauges prior exposure.

📍 SLIDE 2 — Agenda

⏱ Time: 1 minute

Say:

> "Here's our roadmap for the next 60 minutes. Six sections. I'll keep things moving but we'll stop for questions at natural breakpoints. You'll see timings on screen — hold me to them!"

📍 SLIDE 3 — What is a Data Pipeline?

⏱ Time: 5 minutes

Say:

> "Before we jump into ETL and ELT, let's make sure we're all on the same page with what a data pipeline is.

> Think of it like a pizza restaurant. Raw ingredients arrive at the back — that's your Source data. The kitchen preps and cooks them — that's the Transform step. The finished pizza goes into the fridge or directly to the table — that's Store and Consume.
>
> A data pipeline does the exact same thing but with data instead of pizzas. Raw data comes in from somewhere, gets processed, gets stored, and then business teams use it."

Draw on whiteboard: Source → [arrow] → Process → [arrow] → Store → [arrow] → Consume

Say:

> "Now — ETL and ELT are two different philosophies about HOW and WHERE the processing step happens in this pipeline. That's what we're going to explore."

Pause for questions (30 seconds).

📍 SLIDE 4 — ETL: What is it?

⏱ **Time: 5 minutes**

Say:

> "ETL stands for Extract, Transform, Load. Notice the order — Transform comes BEFORE Load. This is the traditional, older approach.
>
> In ETL, you pull your data from sources, you clean and reshape it in a SEPARATE staging area — a dedicated server or tool — and ONLY THEN do you load the clean data into your destination system.

> Think of it like a water filtration plant. You don't dump dirty river water straight into people's homes. You filter it first, THEN distribute it. ETL is that filtration step happening before the data reaches the warehouse."

Point to each letter on the slide: - **E — Extract:** "Pull from databases, CSVs, APIs. Could be 10 sources at once." - **T — Transform:** "This happens on a SEPARATE staging server. Clean, join, validate, enrich." - **L — Load:** "Now the CLEAN data goes into the warehouse — ready to use immediately."

Key emphasis: *"The transformation is happening OUTSIDE the destination. This is crucial."*

📍 SLIDE 5 — ETL Real-World Example (Bank)

⏱ **Time: 4 minutes**

Say:

> "Let's make this concrete. Imagine you work at HDFC Bank as a data engineer.
>
> Every night at 11 PM, your ETL job wakes up. It goes and grabs transaction data from Oracle, account data from MySQL, and KYC information from a CSV file the compliance team exported.
>
> Before any of this touches the warehouse, it goes to a staging server running Informatica. Here, dummy accounts are removed, currencies are all converted to INR, loan utilization percentages are calculated, high-risk accounts are flagged.
>
> THEN — and only then — the perfectly clean, structured data gets loaded into the Teradata Data Warehouse.
>
> Next morning, the risk management team opens their Power BI dashboard and sees clean, accurate data. They have no idea about the overnight transformation marathon that made this possible."

Ask learners: *"Why do you think the bank doesn't process this in real-time? What might go wrong?"* (Elicit: volume, risk of errors, regulatory reasons.)

📍 SLIDE 6 — ELT: What is it?

⏱ **Time: 5 minutes**

Say:

> "Now let's meet ELT — Extract, Load, Transform. Spot the difference: Load comes BEFORE Transform now.
>
> In ELT, you still Extract data from sources. But then you immediately LOAD the raw, unprocessed data directly into your destination — usually a cloud data warehouse like BigQuery or Snowflake.
>
> The transformation happens INSIDE the warehouse using its own compute power.
>
> Why is this possible now when it wasn't before? Because cloud warehouses are MASSIVELY powerful and cheap. In the old days, loading 10 terabytes of raw data into your warehouse and running SQL transformations on it would be cost-prohibitive. Today with BigQuery or Snowflake, it's completely normal.
>
> ELT is the modern cloud-native approach."

Analogy to use:

> "Think of it like a raw food delivery service. You get everything raw and you cook it in your own kitchen using your powerful appliances. ETL is like getting meal-kit delivery where everything is pre-chopped and prepped before it reaches you."

📍 SLIDE 7 — ELT Real-World Example (E-Commerce)

⏱ **Time: 4 minutes**

Say:

> "Let's use Flipkart as our example.
>
> Every few minutes, Fivetran (an ELT ingestion tool) is pulling clickstream data from the website, order data from microservices, and customer profiles from Salesforce. This data lands in Google BigQuery as raw tables — immediately. No waiting.
>
> Hundreds of gigabytes of raw data sitting in BigQuery. Now, dbt — which stands for data build tool — runs SQL models. These SQL transforms run INSIDE BigQuery, using BigQuery's own compute power. It calculates conversion rates, builds customer cohorts, analyzes product performance.
>
> An analyst can query the raw data OR the transformed data. They have access to everything."

Ask: *"What's the big advantage here compared to the bank ETL example?"* (Elicit: raw data available, analysts can explore freely, much faster to load.)

📍 SLIDE 8 — ETL vs ELT Head to Head

⏱ **Time: 5 minutes**

Say:

> "Let's put them side by side. I want you to look at this table and pick out the three most important differences."

Give learners 60 seconds to read silently.

Then highlight:

> "Three things I want you to remember:
>
> Number 1 — WHERE transformation happens. ETL: outside in a staging server. ELT: inside the warehouse.
>
> Number 2 — Raw data access. In ETL, raw data is gone once transformed. In ELT, raw data is always there.
>
> Number 3 — Scale. ETL was designed for megabytes to gigabytes. ELT was designed for terabytes to petabytes."

Write on whiteboard: ETL = Transform first, Load clean | ELT = Load raw, Transform later

---

📍 SLIDE 9 — When to Use ETL vs ELT

⏱ Time: 3 minutes

Say:

> "So which one should you use? The answer is: it depends on your context.
>
> If you're in banking, insurance, or healthcare — sectors with strict regulatory requirements where data must be clean before it touches any system — use ETL.
>
> If you're in e-commerce, SaaS, or tech with terabytes of data and you need analysts to explore freely — use ELT on the cloud.
>
> Many mature companies actually use BOTH — ETL for regulated data, ELT for analytics data. Your job as a data engineer is to know when to pick which."

📍 SLIDE 10 — Batch vs Real-Time Processing

⏱ **Time: 6 minutes**

Say:

> "Now we shift gears slightly. We've been talking about HOW data is processed. Now let's talk about WHEN — and that's the Batch vs Real-Time distinction.
>
> Batch processing means you collect data over a period of time and then process it all at once, usually on a schedule — hourly, nightly, weekly.
>
> Real-time (or stream) processing means you process each event the moment it arrives. Continuously."

Analogy time:

> "Batch is like doing your laundry once a week. You collect dirty clothes all week, then run one big wash cycle. Efficient, but your shirt from Monday won't be clean until Sunday.
>
> Real-time is like having a washing machine that cleans each garment the moment you take it off. Instant, but technically much more complex!"

For Batch:

> "Your salary arrives on the 1st of every month. The payroll system ran a batch job on the last day. Not real-time — no one expects it instantly. Batch is PERFECT here."

For Real-Time:

> "You swipe your credit card at a store. Within 2 seconds you get an SMS confirmation AND the bank's fraud detection system has already checked if this is suspicious. That's real-time. If the bank processed fraud detection in a nightly batch, fraudsters would have 8 hours to drain your account."

Pause and ask: *"Can someone give me another example of where real-time is absolutely necessary versus where batch is fine?"*

📍 SLIDE 11 — Batch vs Real-Time Comparison Table

⏱ **Time: 3 minutes**

Say:

> "Quick comparison table. The key tradeoff is always latency vs complexity.
>
> Batch is simpler, cheaper, but data is stale. Real-time is complex and costs more in compute, but you get instant insights.
>
> Look at the tools row — Spark and Airflow for batch. Kafka and Flink for real-time. We'll see these again in a moment."

📍 SLIDE 12 — Industry Tools & Platforms

⏱ **Time: 6 minutes**

Say:

> "Alright — let's talk tools. The ecosystem here is HUGE. I'm not going to ask you to memorize all of these today, but you should be able to recognize categories and name 2-3 tools per category.
>
> Let me walk through each box:"

Walk through each category:

ETL Tools:

> "Informatica PowerCenter — the enterprise giant. Been around for decades. Used heavily in banks and telecoms. SSIS is Microsoft's ETL tool, part of SQL Server. Talend is open-source and popular in Europe."

ELT / Ingestion:

> "Fivetran and Airbyte are connectors — they extract data from 200+ sources and load it into your warehouse automatically. You don't write extraction code. dbt is the transform layer — you write SQL, dbt runs it inside your warehouse."

Data Warehouses:

> "BigQuery is Google's serverless warehouse. Snowflake is cloud-agnostic and extremely popular. Redshift is Amazon's answer. These are where ELT happens."

Stream / Real-Time:

> "Apache Kafka is the most important tool here — it's a message queue that handles millions of events per second. Think of it as the highway for real-time data. Flink and Spark Streaming process events off that highway."

Orchestration:

> "Airflow is the most popular — it schedules and monitors all your jobs. If your ETL job fails at 2 AM, Airflow sends an alert and retries it."

Data Lakes:

> "S3, Azure Data Lake — object storage for raw data before it goes into a warehouse. Often the 'L' in ELT lands here first."

📍 SLIDE 13 — Modern Data Stack Diagram

⏱ Time: 3 minutes

Say:

> "This is how all these tools fit together in a real company. This is called the Modern Data Stack.
>
> Data flows left to right: Sources → Fivetran extracts and loads into BigQuery/Snowflake → dbt transforms inside the warehouse → Dashboards and ML models consume it.
>
> Below, you have the real-time path: Events → Kafka → Spark Streaming/Flink → Real-time dashboards.
>
> Airflow orchestrates ALL of this — scheduling, monitoring, alerting.

> This exact stack — or something very close to it — is used by companies like Zomato, Razorpay, and most data-mature Indian tech companies today."

📍 SLIDE 14 — Key Takeaways

⏱ **Time: 3 minutes**

Say:

> "Six things I want you to walk out of here with today. Let's read them together..."

Read each one and briefly reinforce: 1. ETL = transform first, outside the warehouse. Legacy/compliance. 2. ELT = load raw first, transform inside. Modern/cloud. 3. Choose based on data volume, governance needs, and infrastructure. 4. Batch = scheduled, high volume, delayed. Real-time = instant, event-driven. 5. Modern stack = Fivetran + BigQuery/Snowflake + dbt + Airflow. 6. Kafka for real-time, Spark/Flink to process streams.

📍 SLIDE 15 — Q&A + Discussion

⏱ **Time: 5 minutes**

Discussion questions (pick 1-2):

Question 1:

> "Your bank sends you a fraud alert within 2 seconds. Which processing type?" **Answer:** Real-time / Stream processing. The transaction event triggers immediate processing in the fraud detection system.

Question 2:

> "A hospital generates monthly patient reports. ETL or ELT?" **Answer:** Likely ETL — healthcare has strict compliance (HIPAA-type regulations), data must be clean before loading, typically on-premise systems.

Question 3:

> "Swiggy recommends a restaurant based on your current browsing. Which pattern?" **Answer:** Real-time streaming — your browsing events are streamed to a recommendation engine that responds instantly.

Closing:

> "Next session we move into Data Modeling — Star Schema and Dimensional Modeling. If you want a head start, look up the difference between a fact table and a dimension table. And if you want hands-on ETL practice, that's Week 13 — so keep this session in mind when we get there. Thank you all — take 10 minutes and we'll come back for the quiz!"

🗒️ COMMON LEARNER QUESTIONS & ANSWERS

Q: Can a company use both ETL and ELT at the same time?

> A: Absolutely. Many enterprises use ETL for regulated, sensitive data and ELT for analytics/ML workloads. The pattern chosen depends on the use case, not the company as a whole.

Q: Is real-time processing always better than batch?

> A: No. Real-time is more complex and expensive. If your business doesn't need instant data (e.g., monthly financial reports), batch is simpler, cheaper, and perfectly appropriate.

Q: What is dbt exactly?

> A: dbt (data build tool) is a transformation framework. You write SQL SELECT statements, and dbt turns them into tables/views inside your warehouse. It's the "T" in ELT. It doesn't extract or load — only transforms.

Q: Is Snowflake a tool for ETL or ELT?

> A: Snowflake is a cloud data warehouse — it's the destination. It's primarily used in ELT patterns where raw data is loaded in and dbt/Spark transforms it inside Snowflake.

Q: What is Apache Kafka used for?

> A: Kafka is a distributed event streaming platform — it acts as a high-throughput message queue. Producers (apps, sensors, microservices) publish events to Kafka topics, and consumers (Flink, Spark, etc.) read and process those events in real-time.

⏱️ TIME TRACKER

Section	Target Time	Cumulative
Title + Agenda	3 min	3 min
Data Pipeline Recap	5 min	8 min

Section	Target Time	Cumulative
ETL Concept + Example	9 min	17 min
ELT Concept + Example	9 min	26 min
ETL vs ELT Comparison	8 min	34 min
Batch vs Real-Time	9 min	43 min
Industry Tools + Stack	9 min	52 min
Takeaways + Q&A	8 min	60 min